

Recurrent Neural Networks

Jean-Martin Albert

August 30, 2017

A Bird's Eye View Of Classical Neural Networks

- ▶ A neural network is a fancy way to describe a class of functions $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ parametrized by \mathbb{R}^ℓ .
- ▶ The layers of a NN represent function composition.
- ▶ If f is a function $F(\mathbf{x}, \mathbf{w}) : \mathbb{R}^n \times \mathbb{R}^\ell \rightarrow \mathbb{R}^m$ such that $\mathcal{C} = \{\mathbf{x} \mapsto F(\mathbf{x}, \mathbf{w}) : \mathbf{w} \in \mathbb{R}^\ell\}$.
- ▶ Each vector $\mathbf{w} \in \mathbb{R}^\ell$ is in some sense a code for an element of \mathcal{C} .
- ▶ When we are given a function $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, we try to find which element of \mathcal{C} best represents f .

A Bird's Eye View Of Classical Neural Networks

- ▶ In tensorflow, the function we are trying to approximate is given in the form of a relation $f(\mathbf{x}) = \mathbf{y}$.
- ▶ We provide *placeholders* for the values of \mathbf{x} and \mathbf{y} .
- ▶ The function $F(\mathbf{x}, \mathbf{w})$ is defined by a *computation graph*.
- ▶ For example, here is a linear regression, in which we try to approximate a function $f : \mathbb{R}^{10} \rightarrow \mathbb{R}^{15}$ using a function of the form $F(\mathbf{x}, \mathbf{W}\mathbf{b}) = \mathbf{W}\mathbf{x} + \mathbf{b}$. Here \mathbf{W} is an 15×10 matrix, and $\mathbf{b} \in \mathbb{R}^{15}$

A Bird's Eye View Of Classical Neural Networks

```
x = tf.placeholder(shape=[None, 10])
y = tf.placeholder(shape=[None, 15])
W = tf.Variable(shape=[10, 15])
b = tf.Variable(shape=[1, 15])
F = tf.matmul(x, W) + b
D = tf.mean_square_error(y, F)
```

- ▶ In tensorflow (and also Theano), vectors are represented as rows instead of columns
- ▶ The first element in the *shape* parameter is there for batching, and None indicates that the dimension can vary.
- ▶ Classical neural networks are great for classification problems involving either fixed (finite) sets $f : A \rightarrow B$, where A represents the set to be classified, and B is the set of labels, or functions $f : \mathbb{R}^n \rightarrow B$.

Sequences

- ▶ In many cases, however, the data to be classified is better represented as a (finite) sequence.
- ▶ For example,
 - ▶ sentences are sequences of words,
 - ▶ words are sequences of letters,
 - ▶ movies are sequences of images,
 - ▶ sound files are sequences of amplitudes.
- ▶ It is possible to deal with sequences with ordinary neural networks, but we run into difficulties when we try to model dependency between different elements of a sequence.

Sequences

Let A be a finite set.

- ▶ A finite sequence of elements of A is a tuple (a_1, \dots, a_n) , where $a_i \in A$ for every i . Here the number n is allowed to change.
- ▶ There is a special sequence, the empty sequence, denoted by ϵ .
- ▶ The set of all finite sequences of elements of A is denoted A^* .

There are several types of functions on sequences:

1. A function $f : A \rightarrow B^*$, which given an element of A outputs a sequence of elements of B . This is called a *one-to-many* function.
2. A function $f : A^* \rightarrow B$, which is given a sequence of elements of A and produces a single element of B . This is called a *many-to-one* function.
3. A function $f : A^* \rightarrow B^*$ which is given a sequence of elements of A and outputs a sequence of elements of B . This is called a *many-to-many* function.

Functions on Sequences

- ▶ We can of course define functions $A \rightarrow B^*$, $A^* \rightarrow B$ and $A^* \rightarrow B^*$ directly (they are just sets, after all),
- ▶ it is more interesting and instructive to make use of the structure of A^* and B^* as sets of sequences of A and B , and see how one can use a function $A \rightarrow B$ to define a function $A^* \rightarrow B^*$.
- ▶ Here is the simplest example: if $A = B$, and $f : A \rightarrow A$ is a function, then we can iterate f : $f^0(a) = a$ for every $a \in A$, and $f^{n+1}(a) = f(f^n(a))$. We terminate the iteration at the first value of n for which $f^n(a) = f^{n-1}(a)$ (note that there is no guarantee that this will ever happen).
- ▶ This gives a function $f^\omega : A \rightarrow A^\omega$.

Functions on Sequences

We describe a more general situation.

- ▶ Let S be a (finite) set of *states* with a distinguished element $\perp \in S$, and
- ▶ consider a function $f : A \times S \rightarrow B \times S$.
- ▶ Let a_1, \dots, a_n be a finite sequence of elements of A .
- ▶ We define two sequences b_1, \dots, b_n and s_1, \dots, s_n of elements of B and S respectively as follows:
 1. $b_1, s_1 = f(a_1, \perp)$.
 2. $b_{n+1}, s_{n+1} = f(a_{n+1}, s_n)$.
 3. We define $f^*(a_1 \dots a_n) = b_1 \dots b_n$, with b_1, \dots, b_n defined as above.

Functions on Sequences

An interesting special case of this is when $S = B$. In this case, we can have $f : A \times B \rightarrow B$. We can define

1. $b_1 = f(a_1, \perp)$,
2. $b_{n+1} = f(a_{n+1}, b_n)$.

This is in fact the main type of function we use in our examples later. Here's a note:

- ▶ Defining a function $f : A \rightarrow B \times C$ is exactly the same as defining two function, $f_1 : A \rightarrow B$ and $f_2 : A \rightarrow C$.
- ▶ In the case of $f : A \times S \rightarrow B \times S$, we are defining two function $f_1 : A \times S \rightarrow B$ and $f_2 : A \times S \rightarrow S$.


The former takes an input and a state to produce an outputs. The latter, taking the same state and input, produces a new state.

Recurrent Neural Networks

- ▶ Recurrent neural networks are just a prescription for a class of functions of the form $f : \mathbb{R}^n \times \mathbb{R}^\ell \times \mathbb{R}^m \rightarrow \mathbb{R}^\ell \times \mathbb{R}^m$, where \mathbb{R}^m is the space of parameters.
- ▶ The set $\{0, 1\}^*$ of all finite sequences of 0's and 1's is countable,
- ▶ and \mathbb{R}^ℓ , is a real vector space, is uncountable.
- ▶ Therefore, we can encode every element $w \in \{0, 1\}^*$ as a vector in S . Informally, \mathbb{R}^ℓ is enough to encode any finite state space, and every possible value for the content of the tape of a Turing machine. We get: *Recurrent neural networks are Turing-complete.*
- ▶ This explains why recurrent neural networks seem to be able to produce results that other networks can't.
- ▶ Training a recurrent network is the same as producing a Turing machine.

Recurrent Neural Networks

A common way to represent a recurrent node is to use a box, the inside of which represents the definition of f . The vertical arrows represent the input and output of the node, and the horizontal arrows represent the input and output state of the node.



Recurrent Neural Networks

- ▶ Just like any regular neural network layer, recurrent nodes can be composed.
- ▶ The composition of two recurrent nodes is done by feeding the output of one node into the input of the other, and concatenating their state space.
- ▶ Formally, if $f : \mathbb{R}^n \times \mathbb{R}^\ell \times \mathbb{R}^m \rightarrow \mathbb{R}^k \times \mathbb{R}^\ell$ and $g : \mathbb{R}^k \times \mathbb{R}^\ell \times \mathbb{R}^m \rightarrow \mathbb{R}^q \times \mathbb{R}^\ell$ are two recurrent nodes,

$$(f; g) : \mathbb{R}^n \times \mathbb{R}^\ell \times \mathbb{R}^m \rightarrow \mathbb{R}^q \times \mathbb{R}^\ell \times \mathbb{R}^l$$

defined by

$$(f; g)(\mathbf{x}, \mathbf{s}_1, \mathbf{s}_2) = (\mathbf{y}, \mathbf{s}'_1, \mathbf{s}'_2)$$

where

$$f(\mathbf{x}, \mathbf{s}_1) = \mathbf{x}', \mathbf{s}'_1$$


and

$$g(\mathbf{x}', \mathbf{s}_2) = \mathbf{y}, \mathbf{s}'_2$$

.


Recurrent Neural Networks

Graphically, we can represent this by stacking the boxes representing f and g on top of one another:



Training RNN's

How do we train recurrent neural networks? We use a variant of back propagation, just like a regular neural network. Heuristically, we unroll the recurrent network “infinitely” many times, until it looks like an ordinary very deep neural network. In practice, we only unroll the network a large but finite number of times, and treat it like an ordinary neural network.



The Basic RNN Cell

We begin with the most basic of recurrent cell.

- ▶ Abstractly, a function $f : U \times S \rightarrow V \times S$ can be defined using two functions $u : \mathbb{R}^n \times \mathbb{R}^\ell \rightarrow \mathbb{R}^n$ and $v : \mathbb{R}^n \times \mathbb{R}^\ell \rightarrow \mathbb{R}^\ell$.
- ▶ The former function as providing the output
- ▶ The latter function as a state updating function.
- ▶ The most common definition for u and v for basic recurrent cells is as follows: $u(x, s) = f(A_u x + B_u s)$, and $v(x, s) = \tanh(A_v x + B_v s)$, where
 - ▶ A_u, A_v, B_u and B_v are matrices,
 - ▶ f is a non-linear function

The Basic RNN Cell (code)

In tensorflow, the code looks like:

```
def basic_rnn_cell(i_tensor, s_tensor, o_dim):  
    i_dim = input_tensor.get_shape()[1]  
    s_dim = input_tensor.get_shape()[1]  
    A_u = tf.Variable(shape=[i_dim, o_dim])  
    B_u = tf.Variable(shape=[s_dim, o_dim])  
    A_v = tf.Variable(shape=[i_dim, s_dim])  
    B_v = tf.Variable(shape=[s_dim, s_dim])  
    o_tensor = tf.relu(tf.matmul(i_tensor, A_u) + \  
                        tf.matmul(s_tensor, B_u))  
    ns_tensor = tf.tanh(tf.matmul(i_tensor, A_v) + \  
                        tf.matmul(s_tensor, B_v))  
    return o_tensor, ns_tensor
```


The Long-Short-Term-Memory Cell

- ▶ It is a cell that can remember elements of a sequence it has already seen
- ▶ Its state has two parts: a state and a memory vector;
- ▶ A special *forget gate* controls how much of the memory vector gets forgotten
- ▶ An *input gate* controls how new information gets stored into the memory vector
- ▶ The usual state update and output functions

The Long-Short-Term-Memory Cell

We define:

- ▶ $u : \mathbb{R}^n \times \mathbb{R}^m \times \mathbb{R}^\ell \rightarrow \mathbb{R}^m$ and
- ▶ $v : \mathbb{R}^n \times \mathbb{R}^m \times \mathbb{R}^\ell \rightarrow \mathbb{R}^\ell$

using auxiliary functions F , I , O , and S defined as follows

$$F(x, y) = \sigma(A_F x + B_F y + b_F)$$

$$I(x, y) = \sigma(A_I x + B_I y + b_I)$$

$$O(x, y) = \sigma(A_O x + B_O y + b_O)$$

$$S(x, y) = \tanh(A_O x + B_O y + b_O)$$

The state update is given by

$$v(x, y, s) = F(x, y) \circ s + I(x, y) \circ \sigma_v(A_v x + B_v y + b_v)$$

where \circ denotes pointwise multiplication of vectors, and finally, the output can be defined as

$$u(x, y, s) = O(x, y) \circ \sigma_u(v(x, y, s))$$

The LSTM Cell (code)

```
def lstm_gate(input_tensor, previous_output, op):
    _, N = input_tensor.get_shape()
    _, output_dim = previous_output.get_shape()
    A = tf.Variable(shape=[N, output_dim])
    B = tf.Variable(shape=[output_dim, output_dim])
    b = tf.Variable(shape=[1, output_dim])
    x = tf.matmul(input_tensor, A) + \
        tf.matmul(previous_output, B) + b
    return op(x)

def lstm_cell(input_tensor, output):
    _, output_dim = output.get_shape()
    F = lstm_gate(input_tensor, output, tf.sigmoid)
    I = lstm_gate(input_tensor, output, tf.sigmoid)
    O = lstm_gate(input_tensor, output, tf.sigmoid)
    S = lstm_gate(input_tensor, output, tf.tanh)
    new_state = tf.mul(output, F) + tf.mul(I, S)
    output = tf.mul(O, tf.tanh(new_state))
    return output, new_state
```

The Basic GRU Cell

- ▶ GRU's are a simplification of the LSTM cell.
- ▶ Compared to LSTM the GRU lacks an output gate.
- ▶ Its performance is on par with LSTM cells in most applications.

To define the functions u and v , we use auxiliary functions $U(\text{pdate})$ and $R(\text{reset})$ defined as follows

$$\begin{aligned}U(x, y) &= \sigma(A_U x + B_U y + b_U) \\R(x, y) &= \sigma(A_R x + B_R y + b_R)\end{aligned}$$

The state update (and output) is given by

$$v(x, y, s) = U(x, y) \circ s + (1 - s) \circ \sigma_h(A_v x + B_v(R(x, y) \circ y) + b_v)$$

The GRU Cell (code)

```
def gru_gate(input_tensor, previous_output, port_op):
    _, N = input_tensor.get_shape()
    _, output_dim = previous_output.get_shape()
    A = tf.Variable(shape=[N, output_dim])
    B = tf.Variable(shape=[output_dim, output_dim])
    b = tf.Variable(shape=[output_dim, output_dim])
    x = tf.matmul(input_tensor, A) + \
        tf.matmul(previous_output, B) + b
    return post_op(x)

def gru_cell(input_tensor, output, state):
    U = gru_gate(input_tensor, output, tf.sigmoid)
    R = gru_gate(input_tensor, output, tf.sigmoid)
    O = gru_gate(input, tf.mul(R, output))
    return [tf.mul(R, output) + tf.mul((1-R), O)]*2
```


A many-to-one example

- ▶ As an example of a many-to-one RNN, consider the Buzzometer sentiment analysis tool.
- ▶ On input of a sequence of **characters**, we output one of 4 classes: negative, neutral, positive and irrelevant.
- ▶ For training, the network was unrolled to 256 characters, which is twice the average length of a message in our database.
- ▶ Longer messages were truncated, and shorter messages were padded with 0.

The architecture is:

- ▶ A 3 layer *bi-directional RNN*
- ▶ We keep the last output of the forward and the backward networks, and combine them linearly: $\mathbf{w} = W_1 \mathbf{v}_f + W_2 \mathbf{v}_b$
- ▶ The vector \mathbf{w} is then projected to \mathbb{R}^4 .

A many-to-one example (model architecture)



A many-to-one example (code)

```
model_input = tf.placeholder(shape=[SEQ_LENGTH])
_ = tf.one_hot(model_input, depth=E_DIM, axis=-1)
_ = tf.reshape(_, [-1, SEQ_LENGTH, E_DIM])
fw = multi_layer_rnn(N_LAYERS, STATE_DIM)
bw = multi_layer_rnn(N_LAYERS, STATE_DIM)
OP = tf.nn.bidirectional_dynamic_rnn
output, _ = OP(fw, bw, _, dtype=tf.float32)
fw_output = tf.reshape(output[0][:, -1:],
                        [-1, STATE_DIM])
bw_output = tf.reshape(output[1][:, :1],
                        [-1, STATE_DIM])
f = project(fw_output, E_DIM)
b = project(bw_output, E_DIM)
e = tf.add(f, b)
model_output = project(e, NUM_CLASSES)
prediction = tf.argmax(model_output, 1)
```


A many-to-one example (output)

	Negative	Neutral	Positive	Irrelevant
Negative	291	113	50	127
Neutral	108	113	50	85
Positive	0	0	0	0
Irrelevant	0	0	0	0

	Negative	Neutral	Positive	Irrelevant
Negative	292	52	36	55
Neutral	61	176	21	62
Positive	5	15	65	15
Irrelevant	33	28	14	70

	Negative	Neutral	Positive	Irrelevant
Negative	319	38	11	43
Neutral	20	188	2	20
Positive	10	18	91	10
Irrelevant	27	27	7	169

A many-to-one example (output)


	Negative	Neutral	Positive	Irrelevant
Negative	334	4	4	9
Neutral	3	283	2	7
Positive	1	0	125	2
Irrelevant	1	0	4	221

	Negative	Neutral	Positive	Irrelevant
Negative	321	16	3	20
Neutral	8	286	4	6
Positive	4	2	125	16
Irrelevant	11	9	1	168

A many-to-many example

- ▶ we define an architecture that generates text in the style of a particular author, or body of text.
- ▶ The architecture is very simple, and consists of three stacked GRU cells.
- ▶ This model will also expose some of the challenges in training recurrent networks.

A many-to-many example (model architecture)



A many-to-one example (output)

- ▶ For training we network was unrolled to 30 characters, and the training was done in batches of 32 strings.
- ▶ The main challenge for training is that the strings should continue from batch to batch.
- ▶ For example, consider training a similar network unrolled to 2 characters with a batch size of 3.
- ▶ Consider the sentence The quick brown fox jumps over the lazy dog.
- ▶ The first step is to cut the string into substrings of length 2:
|Th|e_|qu|ic|k_|br|ow|n_|fo|x_|ju|mp|s_|ov|er|
t|he||l|az|y_|do|g|.|.
- ▶ In a normal batching situation, we would then cut this list of strings into chunks of length 3, like so:
|(Th|e_|qu)|(|ic|k_|br)|(|ow|n_|fo)|(|x_|ju|mp)|
(|s_|ov|er)|(_t|he|_|l)|(|az|y_|do)| but this causes a problem.

A many-to-many example (output)

|Th|e_|qu|ic|k_|br|ow|n_|fo|x_|ju|mp|s_|ov|er|_t|he|_|az|y
|do|g_|.

| (Th|e_|qu) | (ic|k_|br) | (ow|n_|fo) | (x
|ju|mp) | (s_|ov|er) | (_t|he|_|) | (az|y_|do) |

- ▶ The first sequence of the first batch is Th, which will leave the network in a certain state s .
- ▶ For this state to be updated properly, the first element of the second batch should be |e_|, and *not* |ic|.
- ▶ We must therefore use a different batching strategy. The string will give us 7 batches in total, so we number the subsequences with a batch index from 1 to 7 in order, starting over at 1 when we run out of indices.

A many-to-one example (output)

1	2	3	4	5	6	7
Th	e_	qu	ic	k_	br	ow
n_	fo	x_	ju	mp	s_	ov
er	_t	he	_l	az	y_	do
g.						

A many-to-many example (code)

```
model_input = tf.placeholder(shape=[None, SEQ_LENGTH])
initial_state = tf.placeholder(shape=[N_LAYERS, S_DIM])
_ = tf.one_hot(model_input, depth=E_DIM, axis=-1)
encode = multi_layer_rnn(N_LAYERS, STATE_DIM)
state_tuple = tuple(tf.unstack(initial_state, axis=0))
OP = tf.nn.dynamic_rnn
output, state = OP(encode, _,
                   dtype=tf.float32,
                   initial_state=state_tuple)
output = tf.reshape(output, [-1, STATE_DIM])
output = project(output, E_DIM)
out = tf.reshape(out, [-1, SEQ_LENGTH])
model_output = tf.nn.softmax(output)
output = tf.argmax(output, 1)
```


A many-to-many example (text generating)

```
def generate_text(length, session=None):
    generated_text = ''
    character = [[ord('␣')]]
    istate = np.zeros([N_LAYERS, 1, STATE_DIM])
    while len(generated_text) < length:
        feed_dict = {model_input: character,
                     initial_state: istate}
        next_char, state = session.run(
            [out, state], feed_dict=feed_dict)
        op = np.random.multinomial
        next_char_id = op(1, next_char.squeeze(), 1)
        next_char_id = next_char_id.argmax()
        next_char_id = next_char_id \
            if chr(next_char_id) in \
                string.printable else ord("␣")
        generated_text += chr(next_char_id)
        character = [[next_char_id]]
        istate = state
    return generated_text
```

A many-to-many example (output)

A many-to-many example (output)

A many-to-many example (output)

A many-to-many example (output)

A many-to-many example (output)

A many-to-many example (output)

A many-to-many example (output)

A many-to-many example (output)

A many-to-many example (output)

A many-to-many example (output)


A many-to-many example (output)

A many-to-many example (output)

Many-to-one-to-many example

- ▶ A recurrent neural network that can sort sequences
- ▶ Two parts: an encoder, and a decoder
- ▶ The encoder encodes sequences into fixed length vectors
- ▶ The decoder transforms this vector into a sorted list of numbers.
- ▶ For simplicity, the model was restricted to sequences of length 32
- ▶ the elements of the sequence were all between 1 and 128

Many-to-one-to-many example



Many-to-one-to-many example (code for encoder)

```
model_input = tf.placeholder('uint8',  
                             shape=[None, SEQ_LENGTH])  
_ = tf.one_hot(model_input, depth=E_DIM, axis=-1)  
_ = tf.reshape(_, [-1, SEQ_LENGTH, E_DIM])  
encode = multi_layer_rnn(N_LAYERS, STATE_DIM)  
OP = tf.nn.dynamic_rnn  
encoded_input, state = OP(encode, _, dtype=tf.float32)  
encoder_output = state
```


Many-to-one-to-many example (code for decoder)

```
with tf.variable_scope('decoder'):
    training_decoder_input = \
        tf.zeros_like(Globals.model_input)
    _ = tf.one_hot(training_decoder_input,
                   depth=E_DIM, axis=-1)
    _ = tf.reshape(_, [-1, SEQ_LENGTH, E_DIM])
    decode = multi_layer_rnn(N_LAYERS, STATE_DIM)
    OP = tf.nn.dynamic_rnn
    decoded_output, state = OP(decode, _,
                               dtype=tf.float32,
                               initial_state=state)
    decoded_output = tf.reshape(decoded_output,
                                [-1, STATE_DIM])
    output = project(decoded_output, E_DIM)
    out = tf.argmax(output, 1)
```

A many-to-one-to-many example (output)

- ▶ 39, 9, 113, 57, 39, 124, 125, 86, 65, 30, 39, 124, 98, 65, 27, 62, 79,
102, 96, 45, 96, 42, 54, 92, 33, 61, 14, 106, 89, 4, 61, 95;
- ▶ 57, 6, 6, 6, 6, 6, 96, 96, 96, 96, 96, 96, 96, 96, 96, 96, 96, 118, 118,
118, 118, 118, 118, 118, 118, 118, 118, 118, 118, 118, 118
- ▶ 22, 31, 58, 43, 96, 83, 95, 65, 17, 48, 65, 122, 94, 109, 113, 49, 74,
61, 39, 127, 3, 14, 107, 91, 55, 2, 108, 7, 119, 60, 32, 21;
- ▶ 14, 14, 14, 17, 17, 6, 6, 6, 96, 96, 96, 96, 96, 96, 96, 96, 118, 118, 118,
118, 118, 118, 118, 118, 118, 118, 118, 118, 118, 118, 118, 118, 118
- ▶ 117, 15, 77, 4, 57, 121, 108, 76, 72, 119, 41, 4, 94, 44, 103, 124,
120, 20, 79, 121, 68, 12, 31, 97, 93, 19, 45, 27, 119, 92, 119, 92;
- ▶ 65, 65, 95, 95, 95, 96, 96, 96, 96, 96, 96, 96, 96, 96, 96, 96, 96, 118,
118, 118, 118, 118, 118, 118, 118, 118, 118, 118, 118, 118, 118, 118

A many-to-one-to-many example (output)

- ▶ 75, 6, 57, 57, 108, 34, 78, 71, 112, 115, 108, 48, 67, 1, 14, 9, 14, 115, 83, 62, 86, 91, 61, 40, 105, 92, 86, 84, 30, 84, 19, 107;
- ▶ 1, 6, 9, 14, 14, 19, 0, 34, 40, 48, 57, 57, 61, 62, 67, 71, 75, 78, 83, 84, 84, 85, 86, 91, 92, 105, 107, 107, 109, 111, 114, 115
- ▶ 59, 108, 56, 66, 76, 38, 100, 61, 47, 79, 102, 15, 24, 92, 8, 26, 126, 43, 5, 90, 41, 2, 60, 85, 2, 104, 86, 40, 35, 47, 61, 91;
- ▶ 2, 2, 5, 8, 15, 24, 26, 35, 38, 40, 41, 42, 47, 47, 56, 59, 60, 61, 61, 66, 76, 79, 84, 86, 89, 91, 92, 100, 102, 104, 108, 126
- ▶ 25, 50, 64, 4, 40, 47, 6, 14, 97, 32, 87, 103, 44, 25, 84, 40, 95, 13, 113, 66, 38, 79, 106, 40, 26, 16, 74, 50, 119, 32, 80, 16;
- ▶ 4, 6, 13, 14, 16, 16, 25, 25, 26, 32, 32, 38, 39, 40, 40, 44, 47, 50, 50, 64, 66, 74, 79, 80, 84, 86, 95, 97, 103, 106, 113, 119

A many-to-one-to-many example (output)

- ▶ 96, 59, 77, 29, 39, 112, 23, 79, 60, 110, 97, 69, 107, 13, 96, 124, 2, 12, 55, 16, 106, 110, 30, 118, 119, 52, 22, 37, 113, 93, 73, 58;
- ▶ 2, 12, 13, 16, 22, 23, 29, 0, 37, 39, 52, 55, 58, 59, 60, 69, 73, 77, 79, 93, 96, 96, 97, 106, 107, 109, 111, 111, 113, 118, 119, 124
- ▶ 44, 117, 104, 116, 111, 115, 58, 79, 8, 9, 1, 39, 112, 43, 64, 31, 126, 12, 36, 10, 93, 87, 40, 5, 108, 92, 11, 75, 113, 104, 64, 109;
- ▶ 1, 5, 8, 9, 10, 11, 12, 0, 36, 39, 40, 0, 44, 58, 64, 64, 75, 79, 0, 92, 93, 104, 104, 108, 109, 111, 113, 113, 115, 116, 117, 126
- ▶ 58, 19, 10, 52, 22, 61, 67, 96, 3, 7, 116, 96, 54, 24, 19, 45, 127, 124, 11, 114, 53, 75, 126, 84, 122, 41, 75, 1, 119, 18, 92, 51;
- ▶ 1, 3, 7, 10, 11, 18, 19, 19, 22, 24, 41, 45, 51, 52, 53, 54, 58, 61, 67, 75, 75, 84, 92, 96, 96, 114, 116, 119, 122, 124, 126, 127

A many-to-one-to-many example (output)

- ▶ 34, 73, 95, 91, 93, 108, 43, 75, 38, 70, 66, 40, 108, 127, 25, 94, 34, 26, 89, 23, 95, 43, 2, 54, 11, 19, 105, 52, 108, 77, 93, 86;
- ▶ 2, 11, 19, 23, 25, 26, 34, 34, 38, 40, 0, 0, 52, 54, 66, 70, 73, 75, 77, 86, 89, 91, 93, 93, 94, 95, 95, 105, 107, 108, 108, 127
- ▶ 95, 106, 119, 69, 40, 41, 28, 114, 12, 1, 106, 87, 117, 78, 54, 37, 110, 24, 9, 114, 107, 87, 33, 76, 5, 90, 29, 14, 96, 109, 1, 3;
- ▶ 1, 1, 3, 5, 9, 12, 14, 24, 28, 29, 33, 37, 40, 41, 54, 69, 76, 78, 86, 0, 90, 95, 96, 106, 106, 107, 109, 110, 114, 114, 117, 119
- ▶ 122, 97, 90, 61, 72, 66, 60, 60, 25, 125, 84, 73, 114, 46, 112, 76, 110, 62, 58, 34, 126, 124, 102, 35, 11, 100, 47, 113, 85, 64, 22, 89;
- ▶ 11, 22, 25, 34, 35, 46, 47, 58, 60, 60, 61, 62, 64, 66, 72, 73, 76, 84, 85, 89, 90, 97, 100, 102, 110, 113, 113, 114, 122, 124, 125, 126

A many-to-one-to-many example (output)

- ▶ 69, 41, 40, 61, 9, 120, 37, 85, 103, 47, 81, 77, 126, 72, 86, 72, 6, 72, 92, 7, 22, 100, 127, 19, 89, 72, 62, 6, 64, 26, 45, 15;
- ▶ 6, 6, 7, 9, 15, 19, 22, 26, 37, 40, 41, 45, 47, 61, 62, 64, 69, 72, 72, 72, 72, 77, 81, 85, 86, 89, 92, 100, 103, 120, 126, 127
- ▶ 95, 40, 116, 120, 68, 75, 24, 114, 39, 81, 52, 37, 76, 10, 15, 105, 79, 15, 44, 89, 115, 5, 26, 11, 85, 53, 85, 114, 39, 29, 34, 16;
- ▶ 5, 10, 11, 15, 15, 16, 24, 26, 29, 34, 37, 39, 39, 40, 44, 52, 53, 68, 75, 76, 79, 81, 85, 85, 89, 95, 105, 114, 114, 115, 116, 120
- ▶ 20, 61, 78, 15, 61, 8, 108, 36, 85, 96, 40, 80, 106, 24, 66, 82, 96, 43, 126, 33, 80, 116, 35, 86, 98, 81, 76, 89, 6, 103, 117, 16;
- ▶ 6, 8, 15, 16, 20, 24, 33, 35, 36, 40, 0, 61, 61, 66, 76, 78, 80, 80, 81, 82, 85, 86, 89, 96, 96, 98, 103, 106, 108, 116, 117, 126